

Programming project 2 — Coordinate descent

Overview

In this project we consider a standard unconstrained optimization problem:

$$\min L(w)$$

where $L(\cdot)$ is some cost function and $w \in \mathbb{R}^d$. In class, we looked at several approaches to solving such problems—such as gradient descent and stochastic gradient descent—under differentiability conditions on $L(w)$. We will now look at a different, and in many ways simpler, approach:

- Initialize w somehow.
- Repeat: pick a coordinate $i \in \{1, 2, \dots, d\}$, and update the value of w_i so as to reduce the loss.

Two questions need to be answered in order to fully specify the updates:

- (a) Which coordinate to choose?
- (b) How to set the new value of w_i ?

Give answers to these questions, and thereby flesh out a coordinate descent method. For (a), your solution should do something more adaptive than merely picking a coordinate at random.

Then implement and test your algorithm on a *logistic regression* problem. First download the **wine** data set that we have frequented alluded to in class:

<https://archive.ics.uci.edu/ml/datasets/Wine>

This contains 178 data points in 13 dimensions, with 3 classes. Just keep the first two classes (with 59 and 71 points, respectively) so as to create a **binary** problem.

What to turn in

On the due date, upload (to **gradescope**) a **typewritten** report containing the following elements (each labeled clearly).

1. *A short, high-level description of your coordinate descent method.*

In particular, you should give a concise description of how you solve problems (a) and (b) above. Do you need the function $L(\cdot)$ to be differentiable (and maybe even have continuous second-order derivatives), or does it work with any cost function?

2. *Convergence.*

Under what conditions do you think your method converges to the optimal loss? There's no need to prove anything: just give a few sentences of brief explanation.

3. *Experimental results.*

(Remember that all training must take place on just classes 1 and 2.)

- Begin by running a standard logistic regression solver (e.g., from `scikit-learn`) on the training set. It should not be regularized: if the solver forces you to do this, just set the regularization constant suitably to make it irrelevant. Make note of the final loss L^* .
- Then, implement your coordinate descent method and run it on this data.
- Finally, compare to a method that chooses coordinates i uniformly at random and then *updates* w_i *using your method* (we'll call this “random-feature coordinate descent”).
- Produce a clearly-labeled graph that shows how the loss of your algorithm's current iterate—that is, $L(w_t)$ —decreases with t ; it should asymptote to L^* . On the same graph, show the corresponding curve for random-feature coordinate descent.

4. *Critical evaluation.*

Do you think there is scope for further improvement in your coordinate descent scheme in (1); if so, how?

5. (Optional) *Sparse coordinate descent.*

Now, suppose we want a k -sparse solution w : that is, one that has at most k nonzero entries.

- Propose a modified version of your method for this task. Assume k is part of the input, along with the data.
- Do you think this method always find the best k -sparse solution when $L(\cdot)$ is convex?
- Try this out on the `wine` data. Make a table of loss values obtained for a few selected values of k .